

Continual Network Learning

Nicola Di Cicco

Politecnico di Milano, Italy

Amir Al Sadi

Università di Bologna, Italy

Chiara Grasselli

Università di Bologna, Italy

Andrea Melis

Università di Bologna, Italy

Gianni Antichi

Politecnico di Milano, Italy

Massimo Tornatore

Politecnico di Milano, Italy

1 INTRODUCTION

Machine Learning (ML) has lately become a prominent research area for the networking community, with applications in a broad range of topics such as traffic classification [17, 35], routing [11, 33], congestion control [5, 22], and traffic forecasting [21, 36]. In particular, in-network ML has become very attractive, as it allows to leverage the expressiveness of ML models at data plane speed [10, 31, 38]. The common denominator between many in-network ML use-cases is to train a model in the control plane using annotated historical data, and then deploy the model in the data plane for near real-time inference [10, 24, 29, 38]. Unfortunately, the training data will eventually become outdated (a phenomenon formally known as “distribution shift” or “concept drift”), causing the deployed ML model to suffer from performance degradation [18, 19]. While the necessity for frequent model updates has already been raised [6], three fundamental questions remain: (1) **When should we update our model?** The answer is (in theory) fairly simple: *continuously*. We should assume that the input patterns observed by a deployed ML model may change at any point in time; (2) **Which data should we select for model updates?** Again, the answer is (in theory) simple: *only the data that is useful for learning new things*. (3) **What should our model learn?** In principle, *everything*. We want a model that dynamically expands its predictive power without forgetting past experiences.

In this poster, we aim to take a step towards designing a solution that answers those questions. We propose combining Active Learning (AL) [26], which enables filtering relevant information from a vast pool of unannotated data, and Continual Learning (CL) [13], which allows us to learn from streaming data *without forgetting past concepts*. The former, implemented in the switch ASIC, allows us to choose *the right amount of information* that shall be mirrored to the control plane, where the model is updated continually. Finally, the new model can be installed back in the data plane.

Implementing this solution is nontrivial, and needs answering the following research questions: (1) how to implement AL-based filtering in the data plane?; (2) how selective should AL be for network learning?; (3) which ML models are most suitable for continual learning of network traffic?; and (4) how to dynamically reconfigure the data plane?




Figure 1: Distribution shift of TCP flow features from a real-world commercial backbone link [1]. The distribution shift of TCP features such as flow duration and inter-arrival time (IAT) causes performance degradation of a ML model trying to predict them.

2 THE CASE FOR CONTINUAL LEARNING

We run some tests on real-world traces to characterize the amount of distribution shift in TCP flow features. We extracted commonly-used features from real-world traces [1] (e.g., flow duration, inter-arrival time (IAT), and packet size statistics). We observed a shift in the flow duration and in the maximum IAT for small time scales (~one hour) and for large time scales (~a year), respectively. To quantify the impact of these shifts, we consider ad-hoc regression tasks¹ where the targets are either the flow duration or the maximum IAT. We observe² that the test error is significantly larger than training, a phenomenon that is imputable to the observed feature drifts (Fig. 1). Indeed, classical ML models will work properly only if the train and test data are approximately independent and identically distributed [8]. As such, practical in-network ML calls for smart, adaptive approaches.

Literature has been active in proposing *efficient means to offload ML schemes on the data plane* [10, 12, 31, 32, 39]. **Why can't we run existing proposals in a loop?** Pure Online Learning approaches 1) assume that every streamed data point is labeled, and 2) do not pay attention about forgetting as long as the model is fit to the current experience. In our proposal, we want not only to learn adaptively, but also to remember (and therefore exploit) everything that was observed in the past. In this way, *our model will not need additional data for re-learning already-observed concepts*.

¹This is because CAIDA traces do not have task-specific class labels.

²For visualization purposes, we focus on the ranges [5, 20]s for flow duration and [0, 1.5]s for inter-arrival time.

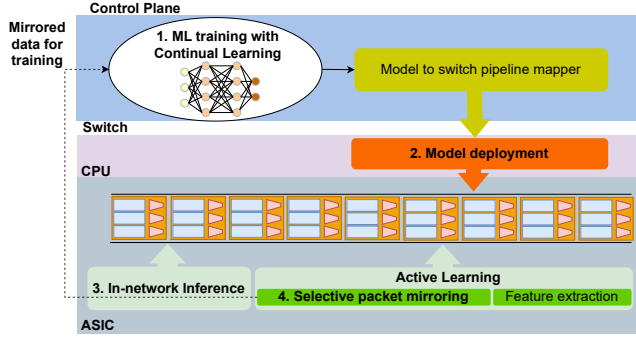


Figure 2: Our approach. The ML model is created (1) and then deployed in the switch ASIC (2), to perform inference at data plane speed (3). Selective mirroring (4) with Active Learning is deployed to keep the ML updated with Continual Learning.

3 OUR APPROACH

Fig. 2 illustrates our idea. We propose to incorporate in a single closed-loop framework the following building blocks:

- (1) Model training: leveraging CL, the ML model updates and expands its knowledge over time.
- (2) Model deployment: the trained ML model is deployed in the data plane (e.g., in a P4-compatible switch).
- (3) In-network inference: the deployed ML model performs inference at data plane speed.
- (4) Selective mirroring: leveraging AL, only the data useful for expanding the knowledge of the deployed model is mirrored to the control plane.

We seek to always run each component: while performing ML inference at data plane speed, our system learns continually by selecting only the most informative data.

As a proof-of-concept experiment, we consider a subset of the CIC2019 dataset for DDoS classification [27]. We consider DDoS classes to represent disjoint learning tasks, which are presented to the model in sequence. For each task, the model must not only discriminate between benign and malicious flows but also place the malicious flows in the right class.

We implement a baseline Continual Random Forest (CRF), consisting of a simple RF augmented with a replay buffer storing the most informative past exemplars. We use the vote count as AL query strategy, selecting only data points whose predictions had less than 90% majority. We retrain after each query, which is computationally very efficient for RFs. We consider Adaptive Random Forest (ARF) as a purely online (but not continual) baseline, which is among the state-of-the-art for Online Learning [14, 19]. In contrast to CRF, ARF assumes that every data point is labeled. We also consider an "oracle" RF trained on the full dataset as an upper-bound on the average performance over all learning tasks.

Fig. 3 shows the performance of CRF and ARF over the sequential tasks, and the percentage of queried labels by CRF




Figure 3: Adaptive vs. Continual Random Forests for class-incremental DDoS classification on CIC2019. At the end of the stream, CRF achieves performance close to an "oracle" while requiring only ~1% of the data.

relative to the full stream size. A purely adaptive learner such as ARF, though able to master individual tasks, quickly forgets past concepts. Instead, our baseline CRF achieves in the end a performance close to the oracle upper-bound, while requiring labeling only ~1% of the observed samples.

4 CHALLENGES

Challenge #1: implementing AL-based filtering in the data plane. While vote count in our baseline CRF is a decent query strategy, information-theoretic quantities [7, 15] are among the state-of-the-art. Their data plane implementation is not trivial, as it would require floating-point arithmetics. Even if not standard, authors in [20] propose a way to implement floating-point arithmetics in P4.

Challenge #2: how selective should AL be. A small selectivity implies a large mirroring overhead, whereas a large selectivity implies a potential information loss. Applying Active Learning to streaming data is, as of today, a novel twist on classical techniques [25]: investigating these trade-offs in a use-case basis opens up interesting research directions.

Challenge #3: choosing the right CL strategy. Our baseline leverages a slowly-growing experience buffer, which may not be desirable. Strategies for maintaining the buffer of fixed size can be investigated [23]. Other solutions, e.g., regularized neural networks, do not require any storage overhead other than the model [16], but are ill-advised for tabular data [28]. Ultimately, the choice depends on the available storage/computational resources and the goodness-of-fit to the characteristics of task-specific data.

Challenge #4: runtime dataplane reconfiguration. Currently, if want to add a new functionality to a switch, we need first to reroute the traffic of that switch, flush a new image in its ASIC and then restore the original traffic policy configuration. This process can lead to dramatic consequences if performed carelessly [2]. Programming the switch at run-time is possible [37], but not for RMT [9], the common commercial devices architecture [3, 4]. Researchers have also explored means to enable isolations between offloaded programs [30, 34, 40], which we will investigate to isolate the Active Learning processing and the rest of the pipeline.

REFERENCES

- [1] [n. d.]. The CAIDA UCSD Anonymized Internet Traces - 2018-2019. ([n. d.]). https://www.caida.org/catalog/datasets/passive_dataset
- [2] [n. d.]. Facebook Outage on October 4th 2021. <https://engineering.fb.com/2021/10/04/networking-traffic/outage/>. ([n. d.]).
- [3] [n. d.]. Naples DSC-100 Distributed Services Card. https://pensando.io/assets/documents/Naples_100_ProductBrief-10-2019.pdf. ([n. d.]).
- [4] [n. d.]. Tofino: P4-programmable Ethernet switch ASIC. <https://www.intel.com/content/www/us/en/products/network-io/programmable-ethernet-switch/tofino-series/tofino.html>. ([n. d.]).
- [5] Soheil Abbasloo, Chen-Yu Yen, and H. Jonathan Chao. 2020. Classic Meets Modern: A Pragmatic Learning-Based Congestion Control for the Internet (SIGCOMM '20). Association for Computing Machinery, New York, NY, USA, 632–647. <https://doi.org/10.1145/3387514.3405892>
- [6] Giovanni Apruzzese, Pavel Laskov, and Johannes Schneider. 2023. SoK: Pragmatic Assessment of Machine Learning for Network Intrusion Detection. In *Proc. IEEE European Symposium on Security and Privacy (EuroSP)*.
- [7] Freddie Bickford Smith, Andreas Kirsch, Sebastian Farquhar, Yarin Gal, Adam Foster, and Tom Rainforth. 2023. Prediction-Oriented Bayesian Active Learning. In *Proceedings of The 26th International Conference on Artificial Intelligence and Statistics (Proceedings of Machine Learning Research)*, Francisco Ruiz, Jennifer Dy, and Jan-Willem van de Meent (Eds.), Vol. 206. PMLR, 7331–7348. <https://proceedings.mlr.press/v206/bickfordsmith23a.html>
- [8] Christopher M. Bishop. 2006. *Pattern Recognition and Machine Learning (Information Science and Statistics)*. Springer-Verlag, Berlin, Heidelberg.
- [9] Pat Bosshart, Glen Gibb, Hun-Seok Kim, George Varghese, Nick McKeown, Martin Izzard, Fernando Mujica, and Mark Horowitz. 2013. Forwarding Metamorphosis: Fast Programmable Match-Action Processing in Hardware for SDN. In *Special Interest Group on Data Communication (SIGCOMM)*. ACM.
- [10] Coralie Busse-Grawitz, Roland Meier, Alexander Dietmüller, Tobias Bühler, and Laurent Vanbever. 2019. pforest: In-network inference with random forests. *arXiv preprint arXiv:1909.05680* (2019).
- [11] Brian Chang, Aditya Akella, Loris D'Antoni, and Kausik Subramanian. 2023. Learned Load Balancing. In *Proceedings of the 24th International Conference on Distributed Computing and Networking (ICDCN '23)*. Association for Computing Machinery, New York, NY, USA, 177–187. <https://doi.org/10.1145/3571306.3571403>
- [12] Bruno Coelho and Alberto Schaeffer-Filho. 2022. BACKORDERS: using random forests to detect DDoS attacks in programmable data planes. In *Proceedings of the 5th International Workshop on P4 in Europe*. 1–7.
- [13] M. Delange, R. Aljundi, M. Masana, S. Parisot, X. Jia, A. Leonardis, G. Slabaugh, and T. Tuytelaars. 2021. A continual learning survey: Defying forgetting in classification tasks. *IEEE Transactions on Pattern Analysis and Machine Intelligence* (2021), 1–1. <https://doi.org/10.1109/TPAMI.2021.3057446>
- [14] Heitor M. Gomes, Albert Bifet, Jesse Read, Jean Paul Barddal, Fabrício Enembreck, Bernhard Pfahringer, Geoff Holmes, and Tael Abdesslem. 2017. Adaptive random forests for evolving data stream classification. *Machine Learning* 106, 9 (Oct 2017), 1469–1495. <https://doi.org/10.1007/s10994-017-5642-8>
- [15] Neil Houlsby, Ferenc Huszár, Zoubin Ghahramani, and Máté Lengyel. 2011. Bayesian Active Learning for Classification and Preference Learning. (2011). [arXiv:stat.ML/1112.5745](https://arxiv.org/abs/1112.5745)
- [16] James Kirkpatrick, Razvan Pascanu, Neil Rabinowitz, Joel Veness, Guillaume Desjardins, Andrei A. Rusu, Kieran Milan, John Quan, Tiago Ramalho, Agnieszka Grabska-Barwinska, Demis Hassabis, Claudia Clopath, Dharshan Kumaran, and Raia Hadsell. 2017. Overcoming catastrophic forgetting in neural networks. *Proceedings of the National Academy of Sciences* 114, 13 (2017), 3521–3526. <https://doi.org/10.1073/pnas.1611835114> [arXiv:https://www.pnas.org/doi/pdf/10.1073/pnas.1611835114](https://www.pnas.org/doi/pdf/10.1073/pnas.1611835114)
- [17] Kang-Hee Lee, Seung-Hun Lee, and Hyun-Chul Kim. 2020. Traffic classification using deep learning: being highly accurate is not enough. In *Proceedings of the SIGCOMM'20 Poster and Demo Sessions*. 1–2.
- [18] Ankur Mallick, Kevin Hsieh, Behnaz Arzani, and Gauri Joshi. 2022. Matchmaker: Data Drift Mitigation in Machine Learning for Large-Scale Systems. In *Proceedings of Machine Learning and Systems*, D. Marculescu, Y. Chi, and C. Wu (Eds.), Vol. 4. 77–94.
- [19] Pavol Mulinka and Pedro Casas. 2018. Adaptive Network Security through Stream Machine Learning. In *Proceedings of the ACM SIGCOMM 2018 Conference on Posters and Demos (SIGCOMM '18)*. Association for Computing Machinery, New York, NY, USA, 4–5. <https://doi.org/10.1145/3234200.3234246>
- [20] Shivam Patel, Rigden Atsatsang, Kenneth M Tichauer, Michael HLS Wang, James B Kowalkowski, and Nik Sultana. 2022. In-network fractional calculations using P4 for scientific computing workloads. In *Proceedings of the 5th International Workshop on P4 in Europe*. 33–38.
- [21] Yu Qiao, Chengxiang Li, Shuzheng Hao, Jun Wu, and Liang Zhang. 2022. Deep or Statistical: An Empirical Study of Traffic Predictions on Multiple Time Scales. In *Proceedings of the SIGCOMM '22 Poster and Demo Sessions (SIGCOMM '22)*. Association for Computing Machinery, New York, NY, USA, 10–12. <https://doi.org/10.1145/3546037.3546048>
- [22] Sudarsanan Rajasekaran, Manya Ghobadi, Gautam Kumar, and Aditya Akella. 2022. Congestion control in machine learning clusters. In *Proceedings of the 21st ACM Workshop on Hot Topics in Networks*. 235–242.
- [23] Sylvestre-Alvise Rebuffi, Alexander Kolesnikov, Georg Sperl, and Christoph H. Lampert. 2017. iCaRL: Incremental Classifier and Representation Learning. In *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*. 5533–5542. <https://doi.org/10.1109/CVPR.2017.587>
- [24] Davide Sanvito, Giuseppe Siracusano, and Roberto Bifulco. 2018. Can the network be the AI accelerator?. In *Proceedings of the 2018 Morning Workshop on In-Network Computing*. 20–25.
- [25] Akanksha Saran, Safoora Yousefi, Akshay Krishnamurthy, John Langford, and Jordan T. Ash. 2023. Streaming Active Learning with Deep Neural Networks. (2023). [arXiv:cs.LG/2303.02535](https://arxiv.org/abs/2303.02535)
- [26] Burr Settles. 2012. *Active Learning*. Springer International Publishing. <https://doi.org/10.1007/978-3-031-01560-1>
- [27] Iman Sharafaldin, Arash Habibi Lashkari, Saqib Hakak, and Ali A. Ghorbani. 2019. Developing Realistic Distributed Denial of Service (DDoS) Attack Dataset and Taxonomy. In *2019 International Carnahan Conference on Security Technology (ICCCST)*. 1–8. <https://doi.org/10.1109/CCST.2019.8888419>
- [28] Ravid Shwartz-Ziv and Amitai Armon. 2021. Tabular Data: Deep Learning is Not All You Need. In *8th ICML Workshop on Automated Machine Learning (AutoML)*. <https://openreview.net/forum?id=vdgtepS1pV>
- [29] Giuseppe Siracusano and Roberto Bifulco. 2018. In-network neural networks. *arXiv preprint arXiv:1801.05731* (2018).
- [30] Radostin Stoyanov and Noa Zilberman. 2020. MTPSA: Multi-Tenant Programmable Switches. In *P4 Workshop in Europe (EuroP4)*. IEEE.
- [31] Tushar Swamy, Alexander Rucker, Muhammad Shahbaz, Ishan Gaur, and Kunle Olukotun. 2022. Taurus: a data plane architecture for per-packet ML. In *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*. 1099–1114.
- [32] Tushar Swamy, Annus Zulfiqar, Luigi Nardi, Muhammad Shahbaz, and Kunle Olukotun. 2023. Homunculus: Auto-Generating Efficient

- Data-Plane ML Pipelines for Datacenter Networks. In *Proceedings of the 28th ACM International Conference on Architectural Support for Programming Languages and Operating Systems, Volume 3*. 329–342.
- [33] Asaf Valadarsky, Michael Schapira, Dafna Shahaf, and Aviv Tamar. 2017. Learning to Route. In *Proceedings of the 16th ACM Workshop on Hot Topics in Networks (HotNets-XVI)*. Association for Computing Machinery, New York, NY, USA, 185–191. <https://doi.org/10.1145/3152434.3152441>
- [34] Tao Wang, Xiangrui Yang, Gianni Antichi, Anirudh Sivaram, and Aurojit Panda. 2022. Isolation Mechanisms for High-Speed Packet-Processing Pipelines. In *Networked Systems Design and Implementation (NSDI)*. USENIX.
- [35] Matthias Wichtlhuber, Eric Strehle, Daniel Kopp, Lars Prepens, Stefan Stegmueller, Alina Rubina, Christoph Dietzel, and Oliver Hohlfeld. 2022. IXP Scrubber: Learning from Blackholing Traffic for ML-Driven DDoS Detection at Scale. In *Proceedings of the ACM SIGCOMM 2022 Conference (SIGCOMM '22)*. Association for Computing Machinery, New York, NY, USA, 707–722. <https://doi.org/10.1145/3544216.3544268>
- [36] Qiong Wu, Xu Chen, Zhi Zhou, Liang Chen, and Junshan Zhang. 2021. Deep reinforcement learning with spatio-temporal traffic forecasting for data-driven base station sleep control. *IEEE/ACM Transactions on Networking* 29, 2 (2021), 935–948.
- [37] Jiarong Xing, Kuo-Feng Hsu, Matty Kadosh, Alan Lo, Yonatan Piasetzky, Arvind Krishnamurthy, and Ang Chen. 2022. Runtime Programmable Switches. In *Networked Systems Design and Implementation (NSDI)*. USENIX.
- [38] Zhaoqi Xiong and Noa Zilberman. 2019. Do Switches Dream of Machine Learning? Toward In-Network Classification. In *Proceedings of the 18th ACM Workshop on Hot Topics in Networks (HotNets '19)*. Association for Computing Machinery, New York, NY, USA, 25–33. <https://doi.org/10.1145/3365609.3365864>
- [39] Changgang Zheng, Zhaoqi Xiong, Thanh T Bui, Siim Kaupmees, Riyad Bensoussane, Antoine Bernabeu, Shay Vargaftik, Yaniv Ben-Itzhak, and Noa Zilberman. 2022. IIsy: Practical in-network classification. *arXiv preprint arXiv:2205.08243* (2022).
- [40] Peng Zheng, Theophilus Benson, and Chengchen Hu. 2018. P4Visor: Lightweight Virtualization and Composition Primitives for Building and Testing Modular Programs. In *Conference on Emerging Networking EXperiments and Technologies (CoNEXT)*. ACM.